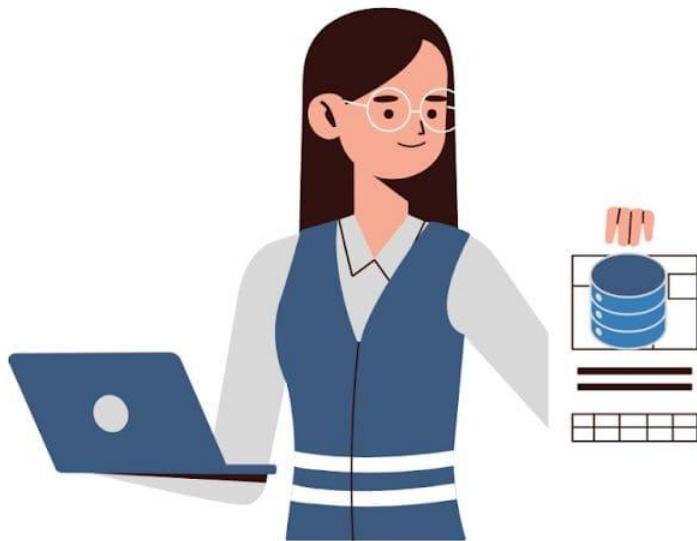


Intro to Databases



Relational Database Models

Lecture 3

Table name: STUDENT

Database name: Ch03_TinyCollege

STU_NUM	STU_LNAME	STU_FNAME	STU_INIT	STU_DOB	STU_HRS	STU_CLASS	STU_GPA	STU_TRANSFER	DEPT_CODE	STU_PHONE	PROF_NUM
321452 Bowser	William	C	12-Feb-1985	42 So	2.84	No BIOL	2134		205		
324257 Smithson	Anne	K.	15-Nov-1991	81 Jr	3.27	Yes CIS	2256		222		
324258 Brewster	Juliette		23-Aug-1979	36 So	2.26	Yes ACCT	2256		228		
324269 Oblonski	Walter	H	16-Sep-1986	66 Jr	3.09	No CIS	2114		222		
324273 Smith	John	D	30-Dec-1968	102 Sr	2.11	Yes ENGL	2231		199		
324274 Katinga	Raphael	P	21-Oct-1989	114 Sr	3.15	No ACCT	2267		228		
324291 Robertson	Gerald	T	08-Apr-1983	120 Sr	3.87	No EDU	2267		311		
324299 Smith	John	B	30-Nov-1996	15 Fr	2.92	No ACCT	2315		230		

What is a TABLE?

- Table – a 2D structure with rows and columns (think of a spreadsheet?)
- Each Row is a record (tuple) with multiple attributes (fields; column names)

TABLE 3.1

CHARACTERISTICS OF A RELATIONAL TABLE

1	A table is perceived as a two-dimensional structure composed of rows and columns.
2	Each table row (tuple) represents a single entity occurrence within the entity set.
3	Each table column represents an attribute, and each column has a distinct name.
4	Each intersection of a row and column represents a single data value.
5	All values in a column must conform to the same data format.
6	Each column has a specific range of values known as the attribute domain .
7	The order of the rows and columns is immaterial to the DBMS.
8	Each table must have an attribute or combination of attributes that uniquely identifies each row.

If you compare to OO Class:
A table is an **array** of Students, with
data fields of stu_num, stu_lname, etc.

Table Attributes (Fields)

- Each data field has a:
 - Distinct name (STU_NUM, STU_LNAME, etc)
 - Data type (just like strongly typed languages, ex: C, Java)
 - Represents a single attribute
- Data Type Options: Numeric, Character (String), Date/Time, Logical (Boolean)
- Valid Data Options: Data fields can have a DOMAIN of possible options
 - GPA can be limited between 0.0 and 4.0 inclusive [0.0, 4.0]
- Order is not important to the user (can be determined later)
- Tables have a PRIMARY KEY, a unique data field used for indexing and protecting from data redundancy (STU_NUM)

Table name: STUDENT

Database name: Ch03_TinyCollege

STU_NUM	STU_LNAME	STU_FNAME	STU_INIT	STU_DOB	STU_HRS	STU_CLASS	STU_GPA	STU_TRANSFER	DEPT_CODE	STU_PHONE	PROF_NUM
321452	Bowser	William	C	12-Feb-1985	42	So	2.84	No	BIOL	2134	205
324257	Smithson	Anne	K	15-Nov-1991	81	Jr	3.27	Yes	CIS	2256	222
324258	Brewer	Juliette		23-Aug-1979	36	So	2.26	Yes	ACCT	2256	228
324269	Oblonski	Walter	H	16-Sep-1986	66	Jr	3.09	No	CIS	2114	222
324273	Smith	John	D	30-Dec-1968	102	Sr	2.11	Yes	ENGL	2231	199
324274	Katinga	Raphael	P	21-Oct-1989	114	Sr	3.15	No	ACCT	2267	228
324291	Robertson	Gerald	T	08-Apr-1983	120	Sr	3.87	No	EDU	2267	311
324299	Smith	John	B	30-Nov-1996	15	Fr	2.92	No	ACCT	2315	230

RE 3.1 STUDENT TABLE ATTRIBUTE VALUES

Table name: STUDENT

Database name: Ch03_TinyCollege

STU_NUM	STU_LNAME	STU_FNAME	STU_INIT	STU_DOB	STU_HRS	STU_CLASS	STU_GPA	STU_TRANSFER	DEPT_CODE	STU_PHONE	PROF_NUM
321452	Bowser	William	C	12-Feb-1985	42	So	2.84	No	BIOL	2134	205
324257	Smithson	Anne	K	15-Nov-1991	81	Jr	3.27	Yes	CIS	2256	222
324258	Brewer	Juliette		23-Aug-1979	36	So	2.26	Yes	ACCT	2256	228
324269	Oblonski	Walter	H	16-Sep-1986	66	Jr	3.09	No	CIS	2114	222
324273	Smith	John	D	30-Dec-1968	102	Sr	2.11	Yes	ENGL	2231	199
324274	Katinga	Raphael	P	21-Oct-1989	114	Sr	3.15	No	ACCT	2267	228
324291	Robertson	Gerald	T	08-Apr-1983	120	Sr	3.87	No	EDU	2267	311
324299	Smith	John	B	30-Nov-1996	15	Fr	2.92	No	ACCT	2315	230

- STU_NUM = Student number
STU_LNAME = Student last name
STU_FNAME = Student first name
STU_INIT = Student middle initial
STU_DOB = Student date of birth
STU_HRS = Credit hours earned
STU_CLASS = Student classification
STU_GPA = Grade point average
STU_TRANSFER = Student transferred from another institution
DEPT_CODE = Department code
STU_PHONE = 4-digit campus phone extension
PROF_NUM = Number of the professor who is the student's advisor

Keys

- Determination – Knowing the value of **1** attribute → determine other attributes.
 - If I know my speed and time, I can DETERMINE my distance ($S * T = D$)
 - In Databases, based on RELATIONSHIPS – Functional Dependence
- Determinate – value that determines other values (dependent)
- Primary Key – Single (usually) data field that is unique to the table (Cannot be null!)
- Composite Keys – Multiple values (key attribute) combine to determine other attribute(s)
- Super Key – Key that can uniquely identify any row (often Primary Key)
- Candidate Key – Minimal Super Key, no unnecessary attributes
 - Used as possible options to determine a primary key

Key	Type
STU_NUM	Primary Key, Candidate Key, Super Key
STU_NUM, STU_LNAME	Composite Key, Super Key
STU_LNAME, STU_FNAME, STU_INIT, STU_PHONE	Composite Key
STU_SSN	Candidate Key, Super Key

Other Keys...

- Foreign Keys – Primary key from one table used as a data field in another.
 - Referential Integrity – Every reference is either valid or NULL
- Secondary Keys
 - Not primary key, but can be used to look up values (Last name, phone number, email?)
 - Could return multiple values, but might be useful anyway

Quick note on NULL:
Null is the absence of data, not a 0 or a space. Can be a problem with functionality (Count, Average, etc)

Table name: PRODUCT

Primary key: PROD_CODE

Foreign key: VEND_CODE

PROD_CODE	PROD_DESCRPT	PROD_PRICE	PROD_ON_HAND	VEND_CODE
001278-AB	Claw hammer	12.95	23	232
123-21UVY	Houselite chain saw, 16-in. bar	189.99	4	235
GER-34256	Sledge hammer, 16-lb. head	18.63	6	231
SRE-657UG	Rat-tail file	2.99	15	232
ZDX3245Q	Steel tape, 12-ft. length	6.79	8	235

Database name: Ch03_SaleCo

Table name: VENDOR

Primary key: VEND_CODE

Foreign key: none

VEND_CODE	VEND_CONTACT	VEND_AREACODE	VEND_PHONE
230	Shelly K. Smithson	608	555-1234
231	James Johnson	615	123-4536
232	Annelise Crystall	608	224-2134
233	Candice Wallace	904	342-6567
234	Arthur Jones	615	123-3324
235	Henry Ortozo	615	899-3425

link

FIGURE 3.2 AN EXAMPLE OF A SIMPLE RELATIONAL DATABASE

Table name: PRODUCT

Primary key: PROD_CODE

Foreign key: VEND_CODE

Database name: Ch03_SaleCo

PROD_CODE	PROD_DESCRIP	PROD_PRICE	PROD_ON_HAND	VEND_CODE
001278-AB	Claw hammer	12.95	23	232
123-21UUY	Houselite chain saw, 16-in. bar	189.99	4	235
QER-34256	Sledge hammer, 16-lb. head	18.63	6	231
SRE-657UG	Rat-tail file	2.99	15	232
ZZX/3245Q	Steel tape, 12-ft. length	6.79	8	235

link

VEND_CODE	VEND_CONTACT	VEND_AREACODE	VEND_PHONE
230	Shelly K. Smithson	608	555-1234
231	James Johnson	615	123-4536
232	Annelise Crystall	608	224-2134
233	Candice Wallace	904	342-6567
234	Arthur Jones	615	123-3324
235	Henry Ortozo	615	899-3425

Table name: VENDOR

Primary key: VEND_CODE

Foreign key: none

Integrity Rules

Entity Integrity –
Primary Keys will be
unique and NOT null

Referential Integrity –
Foreign Keys will either
be null or a valid
primary key in other
table

INTEGRITY RULES	
ENTITY INTEGRITY	DESCRIPTION
Requirement	All primary key entries are unique, and no part of a primary key may be null.
Purpose	Each row will have a unique identity, and foreign key values can properly reference primary key values.
Example	No invoice can have a duplicate number, nor can it be null; in short, all invoices are uniquely identified by their invoice number.
REFERENTIAL INTEGRITY	
Requirement	A foreign key may have either a null entry, as long as it is not a part of its table's primary key, or an entry that matches the primary key value in a table to which it is related (every non-null foreign key value <i>must</i> reference an existing primary key value).
Purpose	It is possible for an attribute <i>not</i> to have a corresponding value, but it will be impossible to have an invalid entry; the enforcement of the referential integrity rule makes it impossible to delete a row in one table whose primary key has mandatory matching foreign key values in another table.
Example	A customer might not yet have an assigned sales representative (number), but it will be impossible to have an invalid sales representative (number).

Relational Algebra

- RelVar – Relational Variable –
 - Container that holds our relation (reference vs. actual data)
 - Can be thought of as a reference to a table or a record-set of data
- Relational Set Operators
 - Closure – Use of relational algebra operators to produce new relations

Select	<ul style="list-style-type: none">• Restrict – Yields all values from 1 table that meet a specific criteria
Project	<ul style="list-style-type: none">• Returns all values for selected attributes from 1 table
Union	<ul style="list-style-type: none">• Combines all rows from 2 tables, excluding duplicate rows. Tables must have same attributes!
Intersect	<ul style="list-style-type: none">• Yields ONLY rows that exists in both tables (2 tables). Tables must have same attributes.
Difference	<ul style="list-style-type: none">• Yields All rows in one table (A) that are NOT in another table (B)
Product	<ul style="list-style-type: none">• Yields all possible pairs of rows from two tables (Cartesian Product)
Join	<ul style="list-style-type: none">• Allows information combined from 2 or more tables (VERY, VERY IMPORTANT!)
Divide	<ul style="list-style-type: none">• Divide two tables: dividend must be 2 columns, divisor must be 1 column... return any records in the second column of the first table in which the first column matches all values in the divisor

Joining 2 (or more) tables

the real power behind relational databases

- **Natural Join** – Product → Select → Project :
 - Gets all rows from both tables where the ‘on’ criteria are equal
 - “SELECT * from A JOIN B on A.stu_num = B.stu_num” : Get all fields from both tables where the student number is equal
- **EquiJoin** – Link tables with an equality condition
- **ThetaJoin** – Link tables with a inequality comparator ($>$, $<$, \leq , \geq)
- **Outer Join** (Left or Right)
 - Join that matches ALL values from one table, and matching values from the other table (depending on which side of the equation)
 - Select * from A LEFT JOIN B on A.stu_num = B.stu_num : Get all fields from A, and only those from B where the students numbers are the same.
 - Select * from A RIGHT JOIN B on A.stu_num = B.stu_num: Get all fields from B and only those from A where the student numbers are the same
 - Non-matching fields will return as null

1:M Relationships in the Relational Database

- 1:M – Easiest to implement and maintain (try to make this!)
 - 1: Primary key in this table – Painter_Num (Painter)
 - M: Foreign key (Painter_Num) in this table (Painting)

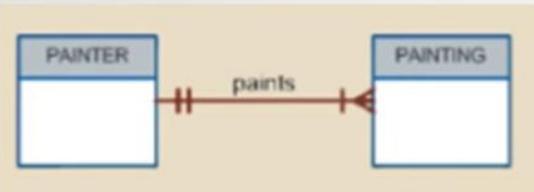


Table name: PAINTER

Primary key: PAINTER_NUM

Foreign key: none

Database name: Ch03_Museum

PAINTER_NUM	PAINTER_LNAME	PAINTER_FNAME	PAINTER_INITIAL
123	Ross	Georgette	P
126	Itero	Julio	O

Table name: PAINTING

Primary key: PAINTING_NUM

Foreign key: PAINTER_NUM

PAINTING_NUM	PAINTING_TITLE	PAINTER_NUM
1338	Dawn Thunder	123
1339	Vanilla Roses To Nowhere	123
1340	Tired Flounders	126
1341	Hasty Exit	123
1342	Plastic Paradise	126

The 1:M Relationship

THE IMPLEMENTED 1:M RELATIONSHIP BETWEEN PAINTER AND PAINTING

Table name: PAINTER

Primary key: PAINTER_NUM

Foreign key: none

	PAINTER_NUM	PAINTER_LNAME	PAINTER_FNAME	PAINTER_INITIAL
	123	Ross	Georgette	P
	126	Itero	Julio	G

Database

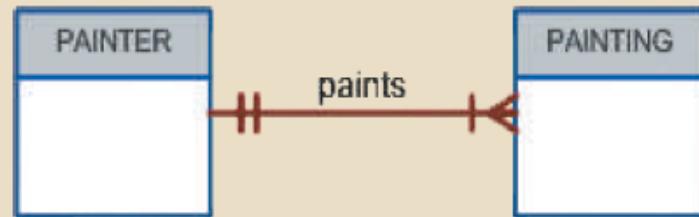


Table name: PAINTING

Primary key: PAINTING_NUM

Foreign key: PAINTER_NUM

PAINTING_NUM	PAINTING_TITLE	PAINTER_NUM
1338	Dawn Thunder	123
1339	Vanilla Roses To Nowhere	123
1340	Tired Flounders	126
1341	Hasty Exit	123
1342	Plastic Paradise	126

The 1:M Relationship

20 THE IMPLEMENTED 1:M RELATIONSHIP BETWEEN COURSE AND CLASS

1:M RELATIONSHIP BETWEEN COURSE AND CLASS

Table name: COURSE

Primary key: CRS_CODE

Foreign key: none

CRS_CODE	DEPT_CODE	CRS_DESCRIPTION	CRS_CREDIT
ACCT-211	ACCT	Accounting I	3
ACCT-212	ACCT	Accounting II	3
CIS-220	CIS	Intro. to Microcomputing	3
CIS-420	CIS	Database Design and Implementation	4
QM-261	CIS	Intro. to Statistics	3
QM-362	CIS	Statistical Applications	4

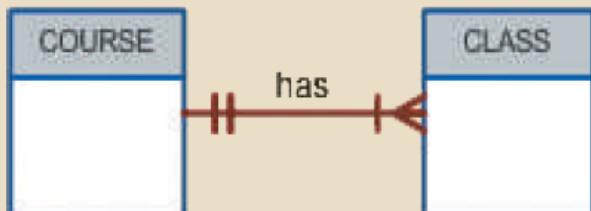
Database name: Ch03_TinyDB

Table name: CLASS

Primary key: CLASS_CODE

Foreign key: CRS_CODE

CLASS_CODE	CRS_CODE	CLASS_SECTION	CLASS_TIME	CLASS_ROOM	PROF_NUM
10012	ACCT-211	1	MWF 8:00-8:50 a.m.	BUS311	105
10013	ACCT-211	2	MWF 9:00-9:50 a.m.	BUS200	105
10014	ACCT-211	3	TTh 2:30-3:45 p.m.	BUS252	342
10015	ACCT-212	1	MWF 10:00-10:50 a.m.	BUS311	301
10016	ACCT-212	2	Th 6:00-8:40 p.m.	BUS252	301
10017	CIS-220	1	MWF 9:00-9:50 a.m.	KLR209	228
10018	CIS-220	2	MWF 9:00-9:50 a.m.	KLR211	114
10019	CIS-220	3	MWF 10:00-10:50 a.m.	KLR209	228
10020	CIS-420	1	W 6:00-8:40 p.m.	KLR209	162
10021	QM-261	1	MWF 8:00-8:50 a.m.	KLR200	114
10022	QM-261	2	TTh 1:00-2:15 p.m.	KLR200	114
10023	QM-362	1	MWF 11:00-11:50 a.m.	KLR200	162
10024	QM-362	2	TTh 2:30-3:45 p.m.	KLR200	162





1:1 Relationships

- Should be rare (try to avoid)
- Should be optional (otherwise, you'd have the data in the same table?)
- Add primary key from first table as the foreign key in the second table.
- Add the primary key from the second table as the foreign key in the first table.

Table name: PROFESSOR
Primary key: EMP_NUM
Foreign key: DEPT_CODE

EMP_NUM	DEPT_CODE	PROF_OFFICE	PROF_EXTENSION	PROF_HIGH_DEGREE
103 HIST	DRE 156	6783		Ph.D.
104 ENG	DRE 102	5561		MA
105 ACCT	KLR 229D	8665		Ph.D.
106 MKTMGT	KLR 126	3099		Ph.D.
110 BIOL	AAK 160	3412		Ph.D.
114 ACCT	KLR 211	4436		Ph.D.
155 MATH	AAK 201	4440		Ph.D.
160 ENG	DRE 102	2248		Ph.D.
162 CIS	KLR 203E	2359		Ph.D.
191 MKTMGT	KLR 409B	4016		DBA
195 PSYCH	AAK 297	3550		Ph.D.
209 CIS	KLR 333	3421		Ph.D.
228 CIS	KLR 300	3000		Ph.D.
297 MATH	AAK 194	1145		Ph.D.
299 ECONFIN	KLR 294	2651		Ph.D.
301 ACCT	KLR 244	4683		Ph.D.
336 ENG	DRE 208	2000		Ph.D.
342 SOC	BBO 208	5514		Ph.D.
387 BIOL	AAK 230	8665		Ph.D.
401 HIST	DRE 156	6783		MA
425 ECONFIN	KLR 284	2651		MBA
435 ART	BBO 185	2278		Ph.D.

Database name: Ch03_TinyCollege

The 1:M DEPARTMENT employs PROFESSOR relationship is implemented through the placement of the DEPT_CODE foreign key in the PROFESSOR table.

Table name: DEPARTMENT
Primary key: DEPT_CODE
Foreign key: EMP_NUM

DEPT_CODE	DEPT_NAME	SCHOOL_CODE	EMP_NUM	DEPT_ADDRESS	DEPT_EXTENSION
ACCT	Accounting	BUS	114	KLR 211, Box 52	3119
ART	Fine Arts	A&SO	435	BBO 185, Box 128	2278
BIOL	Biology	A&SO	387	AAK 230, Box 415	4117
CIS	Computer Info. Systems	BUS	209	KLR 333, Box 56	3245
ECON/FIN	Economics/Finance	BUS	299	KLR 284, Box 63	3126
ENG	English	A&SO	160	DRE 102, Box 223	1004
HIST	History	A&SO	103	DRE 156, Box 264	1067
MATH	Mathematics	A&SO	297	AAK 194, Box 422	4234
MKTMGT	Marketing/Management	BUS	106	KLR 126, Box 55	3342
PSYCH	Psychology	A&SO	195	AAK 297, Box 438	4110
SOC	Sociology	A&SO	342	BBO 208, Box 132	2008

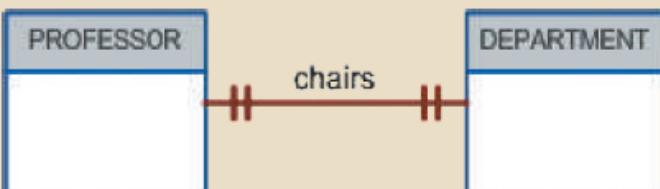
The 1:1 PROFESSOR chairs DEPARTMENT relationship is implemented through the placement of the EMP_NUM foreign key in the DEPARTMENT table.

FIGURE 3.22 THE IMPLEMENTED 1:1 RELATIONSHIP BETWEEN PROFESSOR AND DEPARTMENT

Table name: PROFESSOR
 Primary key: EMP_NUM
 Foreign key: DEPT_CODE

EMP_NUM	DEPT_CODE	PROF_OFFICE	PROF_EXTENSION	PROF_HIGH_DEGREE
103	HIST	DRE 156	6763	Ph.D.
104	ENG	DRE 102	5561	M.A.
105	ACCT	KLR 229D	8665	Ph.D.
106	MKT/MGT	KLR 126	3899	Ph.D.
110	BIO/L	AAK 160	3412	Ph.D.
114	ACCT	KLR 211	4436	Ph.D.
155	MATH	AAK 201	4440	Ph.D.
160	ENG	DRE 102	2248	Ph.D.
162	CIS	KLR 203E	2359	Ph.D.
191	MKT/MGT	KLR 409B	4016	DBA
195	PSYCH	AAK 297	3550	Ph.D.
209	CIS	KLR 333	3421	Ph.D.
228	CIS	KLR 300	3000	Ph.D.
297	MATH	AAK 194	1145	Ph.D.
299	ECON/FIN	KLR 284	2851	Ph.D.
301	ACCT	KLR 244	4583	Ph.D.
335	ENG	DRE 208	2000	Ph.D.
342	SOC	BBG 206	5514	Ph.D.
387	BIO/L	AAK 230	8665	Ph.D.
401	HIST	DRE 156	6783	M.A.
425	ECON/FIN	KLR 284	2851	DBA
435	ART	BBG 185	2278	Ph.D.

Database name: Ch03_TinyCollege



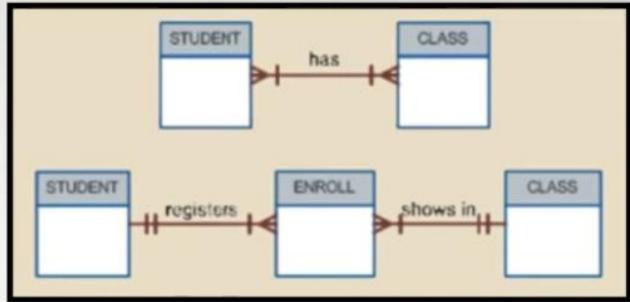
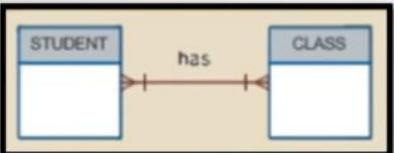
The 1:M DEPARTMENT employs PROFESSOR relationship is implemented through the placement of the DEPT_CODE foreign key in the PROFESSOR table.

Table name: DEPARTMENT
 Primary key: DEPT_CODE
 Foreign key: EMP_NUM

DEPT_CODE	DEPT_NAME	SCHOOL_CODE	EMP_NUM	DEPT_ADDRESS	DEPT_EXTENSION
ACCT	Accounting	EUS	114	KLR 211, Box 52	3119
ART	Fine Arts	A&SCI	435	BBG 185, Box 128	2278
BIO/L	Biology	A&SCI	387	AAK 230, Box 415	4117
CIS	Computer Info. Systems	EUS	209	KLR 333, Box 56	3245
ECON/FIN	Economics/Finance	EUS	299	KLR 284, Box 63	3126
ENG	English	A&SCI	160	DRE 102, Box 223	1004
HIST	History	A&SCI	103	DRE 156, Box 284	1867
MATH	Mathematics	A&SCI	297	AAK 194, Box 422	4234
MKT/MGT	Marketing/Management	EUS	106	KLR 126, Box 55	3342
PSYCH	Psychology	A&SCI	195	AAK 297, Box 436	4110
SOC	Sociology	A&SCI	342	BBG 208, Box 132	2008

The 1:1 PROFESSOR chairs DEPARTMENT relationship is implemented through the placement of the EMP_NUM foreign key in the DEPARTMENT table.

M:N Relationships



- M:N – Can't be done directly... must be implemented into 1:M / N:1 relationships
- Create middle table (linking table) to hold middle relationship

Database name: Ch03_CollegeTry

Table name: STUDENT
Primary key: STU_NUM
Foreign key: none

STU_NUM	STU_LNAME	CLASS_CODE
321452	Bowser	10014
321452	Bowser	10018
321452	Bowser	10021
324257	Smithson	10014
324257	Smithson	10018
324257	Smithson	10021

Table name: CLASS
Primary key: CLASS_CODE
Foreign key: STU_NUM

CLASS_CODE	STU_NUM	CRS_CODE	CLASS_SECTION	CLASS_TIME	CLASS_ROOM	PROF_NUM
10014	321452	ACCT-211	3	TTh 2:30-3:45 p.m.	BUS252	342
10014	324257	ACCT-211	3	TTh 2:30-3:45 p.m.	BUS252	342
10018	321452	CIS-220	2	MWF 9:00-9:50 a.m.	KLR211	114
10018	324257	CIS-220	2	MWF 9:00-9:50 a.m.	KLR211	114
10021	321452	QM-261	1	MWF 8:00-8:50 a.m.	KLR200	114
10021	324257	QM-261	1	MWF 8:00-8:50 a.m.	KLR200	114

- Now you have two 1:M relationships
 - Student : Enroll (1:M)
 - Class : Enroll (1:N)

Database name: Ch03_CollegeTry2

Table name: STUDENT
Primary key: STU_NUM
Foreign key: none

STU_NUM	STU_LNAME
321452	Bowser
324257	Smithson

Table name: ENROLL
Primary key: CLASS_CODE + STU_NUM
Foreign key: CLASS_CODE, STU_NUM

CLASS_CODE	STU_NUM	ENROLL_GRADE
10014	321452	C
10014	324257	B
10018	321452	A
10018	324257	B
10021	321452	C
10021	324257	C

Table name: CLASS
Primary key: CLASS_CODE
Foreign key: CRS_CODE

CLASS_CODE	CRS_CODE	CLASS_SECTION	CLASS_TIME	CLASS_ROOM	PROF_NUM
10014	ACCT-211	3	TTh 2:30-3:45 p.m.	BUS252	342
10018	CIS-220	2	MWF 9:00-9:50 a.m.	KLR211	114
10021	QM-261	1	MWF 8:00-8:50 a.m.	KLR200	114

Data Redundancy - Normalization

- Use of proper foreign keys minimizes redundancy.
- Real test: Does the elimination of an attribute eliminate information
 - Can I get the information through other means?
 - Is the cost of the lookup worth the savings in data redundancy?
 - How is the data used in normal use?

Indexes

- An ordered array of key values and ROW ID values (pointers) to help speed up lookups.
- Helps find information quickly.
- Index key (ex: primary key, but not required) – Index reference point
- Primary keys and foreign keys are often indexed, since you often join on those fields.
- Unique Index – Only 1 row associated with it.
- Indexes are assigned to a single table, but a single table can have many indexes

Codd's Relational Database Rules

TABLE 13.8

DR. CODD'S 12 RELATIONAL DATABASE RULES

RULE	RULE NAME	DESCRIPTION
1	Information	All information in a relational database must be logically represented as column values in rows within tables.
2	Guaranteed access	Every value in a table is guaranteed to be accessible through a combination of table name, primary key value, and column name.
3	Systematic treatment of nulls	Nulls must be represented and treated in a systematic way, independent of data type.
4	Dynamic online catalog based on the relational model	The metadata must be stored and managed as ordinary data—that is, in tables within the database; such data must be available to authorized users using the standard database relational language.
5	Comprehensive data sublanguage	The relational database may support many languages; however, it must support one well-defined, declarative language as well as data definition, view definition, data manipulation (interactive and by program), integrity constraints, authorization, and transaction management (begin, commit, and rollback).
6	View updating	Any view that is theoretically updatable must be updatable through the system.
7	High-level insert, update, and delete	The database must support set-level inserts, updates, and deletes.
8	Physical data independence	Application programs and ad hoc facilities are logically unaffected when physical access methods or storage structures are changed.
9	Logical data independence	Application programs and ad hoc facilities are logically unaffected when changes are made to the table structures that preserve the original table values (changing order of columns or inserting columns).
10	Integrity independence	All relational integrity constraints must be definable in the relational language and stored in the system catalog, not at the application level.
11	Distribution independence	The end users and application programs are unaware of and unaffected by the data location (distributed vs. local databases).
12	Nonsubversion	If the system supports low-level access to the data, users must not be allowed to bypass the integrity rules of the database.
13	Rule zero	All preceding rules are based on the notion that to be considered relational, a database must use its relational facilities exclusively for management.

Summary

- Tables are the basic building blocks of relational databases
- Keys are central to relational tables (including a primary key)
- Tables are linked through common attributes (and foreign keys)
- Primary Operations: Select, Join
 - Secondary Operations: Project, Intersect, Union, Difference, Product, Divide
- Data Dictionary/System Catalog holds metadata about the database
- Good design starts by understanding relationships (1:1, 1:M, M:N)