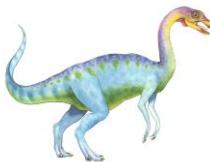


# Operating System 2- CS402

## Lecture 2 Disk scheduling

2025



# Learning Objectives

---

- Learn the process of **disk initialization and formatting.**  
*(Explicitly includes formatting, which is a key part of initialization.)*
- Understand **disk structure and organization.**
- Explore different **techniques for disk scheduling.**



# Disk Management

---

- Computers use disks as the principal **on-line storage** medium for both programs and data.
- Programs are stored on a disk until loaded into memory. They then use the disk as both **the source and destination** of their processing.
- The operating system is responsible for several aspects of disk management. Here we discuss:
  - Disk formatting.
  - Booting from disk.
  - Bad-block recovery.




# Disk Formatting

---

A new magnetic disk is a blank slate: it is just a platter of a magnetic recording material.



1. **Physical Formatting** (Low-Level Formatting)
2. **Partitioning** (Optional but Recommended)
3. **Logical Formatting** (High-Level Formatting)






# Physical Formatting (Low-Level Formatting)

- **Purpose:**
  - Divides the disk into tracks and sectors.
  - Marks bad sectors to avoid data corruption.
- **Performed By:**
  - Manufacturers before shipping the disk.
  - Can be done manually using specialized low-level formatting.
- **Result:** The disk is structured but **not yet ready for data storage**. It still needs a file system.





# Partitioning (Optional but Recommended)

## ■ Purpose:

Partitioning is the process of dividing the hard-disk into one or more regions (one or more groups of cylinders). The regions are called as partitions.


## ■ Types of Partitions:

- **Primary Partition** (Bootable, where OS is installed).
- **Logical Partition** (Used for storing data, applications, etc.).
- **Extended Partition** (cannot store data directly but contain multiple logical partitions).

## ■ Tools to Create Partitions:

- **Windows:** Disk Management (diskpart)
- **Linux:** fdisk, gparted






# Logical Formatting (High-Level Formatting)

## Purpose:

- Creates a **file system** (e.g., NTFS, FAT32, ext4) on the partition to manage data storage.
- Creates a **boot sector**, file allocation tables, and directory structures.
- Makes the disk **usable for storing files**.

## Tools to Perform Logical Formatting:

- Windows:** Right-click disk → Format
- Linux:** mkfs -t ext4 /dev/sda1
- Mac:** Disk Utility



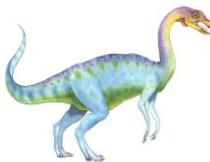


# Bad-Block Recovery

---

- The process of identifying and handling defective blocks (or sectors) on a storage device to prevent data loss and maintain system reliability.
- Bad blocks can develop due to:
  - Manufacturing defects.
  - Aging.
  - Physical damage.





# Bad Blocks Detection Methods

---


- **Error-Correcting Code (ECC):** Storage devices use ECC to detect and correct minor errors in bad sectors (Immediate Error Correction).
- **Surface Scanning:** a low-level disk check that detects bad sectors (physical damage or corruption). Disk utilities (e.g., chkdsk on Windows, badblocks on Linux) scan for unreadable sectors.
- **SMART Monitoring (Self-Monitoring, Analysis, and Reporting Technology):** a built-in monitoring system in HDDs and SSDs that predicts hardware failures. It tracks disk health based on various performance indicators.

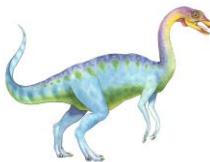


# Sector Data Structure

---

- Low-level formatting fills the disk with a special data structure for each sector.
- The **data structure** for a sector typically consists of:
  - Header.
  - Data area (usually 512 bytes in size).
  - Trailer.
- The header and trailer contain information used by the disk controller, such as a sector number and an **error-correcting code (ECC)**.





# Error-Correcting Code (ECC)

---

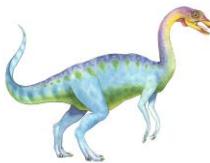
- When the controller writes a sector of data during normal I/O, the **ECC is updated** with a value calculated from all the bytes in the data area.
- When the sector is read the **ECC is recalculated** and compared with the stored value. If the stored and calculated numbers are different, this mismatch indicates that the **data area of the sector has become corrupted** and that the disk sector may be bad.
- The **controller automatically does the ECC processing** whenever a sector is read or written.



# Error-Correcting Code (ECC) (Cont.)

---

- The ECC is an error-correcting code because It contains enough information, if only a few bits of data have been corrupted, to enable the controller to:
  - Identify which bits have changed.
  - Calculate what their correct values should be.
  - Reports a recoverable **soft error**.



# Bootstrapping & System Initialization

---

**Bootstrap Program:** The initial program executed when a computer is powered on or rebooted.


- **Functions of the Bootstrap Program:**
  - Initializes CPU registers, device controllers, and memory.
  - Locates the operating system (OS) kernel on disk.
  - Loads the kernel into memory and starts OS execution.
- **Storage of Bootstrap Program:**
  - Stored in ROM – Non-volatile and immune to viruses.
  - **Limitations** – ROM cannot be easily modified.
- **Solution: Boot loader in ROM**
  - A small bootstrap loader in ROM loads a full bootstrap program from disk.
  - The full bootstrap is stored in boot blocks of a boot disk (system disk).
- **Key Benefit:** Allows easy updates to the bootstrap without hardware changes.



# Boot process in Windows

- Windows divides the hard disk into partitions, with one designated as the **boot partition**, containing the OS and device drivers.
- Master Boot Record (MBR):** The first sector of the hard disk, the MBR, **holds** boot code, partition details, and a flag for the active boot partition.
- Boot Process Steps:**
  - ROM Code Execution:** The system starts by running boot code from ROM.
  - MBR Loading:** The ROM directs the system to load boot code from the MBR.
- Boot Sector Loading:** The system identifies the boot partition, reads its boot sector (the first sector), and continues loading OS components.


This process ensures Windows initializes essential subsystems and system services.





# Disk Addressing

- For efficiently manage data access and retrieval **addressing** is done during Low-level formatting:
  - Disk addressing**: determines how data is located on the disk by the CHS (Cylinder, Head, Sector) scheme that uses three numbers.
  - Logical Block Addressing (LBA)**: creates a 1D array of uniquely numbered logical blocks (smallest unit of transfer).
  - Logical Block Mapping**: translates logical storage structures into physical locations on the disk.






# Disk Addressing

(Cont.)

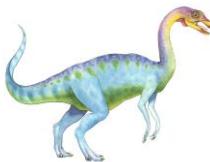
- Disk addressing:
  - Sector 0 → First sector of the first track on the outermost cylinder.
  - Mapping follows track-by-track, cylinder-by-cylinder (outer to inner).
- Logical blocks are **sequentially** mapped to disk sectors.
- Logical to Physical Mapping is straightforward **except**:
  - Bad sectors (unusable areas).
  - **Non-constant** number of sectors per track (due to varying disk geometry).



# HDD Scheduling

---

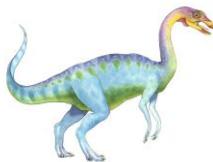
- The operating system is responsible for using hardware efficiently — for the disk drives, this means having a **fast access time** and **large disk bandwidth**.
- **Access time** has two major components:
  - **The seek time** is the time for the disk arm to move the heads to the cylinder containing the desired sector.
  - The **rotational latency** is the additional time for the disk to rotate the desired sector to the disk head.
- **Disk bandwidth** is the total number of bytes transferred, divided by the total time between the first request for service and the completion of the last transfer.



# Disk Scheduling

(Cont.)

- Whenever a process needs I/O to or from the disk, it issues a **system call** to the operating system.
- **The request specifies** several pieces of information:
  - Whether this operation is input or output.
  - What is the disk address for the transfer.
  - What is the memory address for the transfer.
  - What is the number of sectors to be transferred.
- **If the desired disk drive and controller are available**, the request can be serviced **immediately**.
- **If the drive or controller is busy**, any new requests for service will be placed in the queue of **Pending Requests** for that drive.
- When one request is completed, the operating system **chooses** which pending request to service next.





# Disk Scheduling

(Cont.)

- Several **algorithms** exist to schedule the servicing of disk I/O requests.
- The analysis is true for one or many **platters**.
- We illustrate scheduling algorithms with a request queue (0-199).

Queue = 98, 183, 37, 122, 14, 124, 65, 67  
Head pointer at 53






# First-Come, First-Served (FCFS)

---

- The **simplest** form of disk scheduling.
- This algorithm is fair, **but it generally does not provide the fastest service.**
- Consider, for example, a disk queue with requests for I/O to blocks on cylinders  
**98, 183, 37, 122, 14, 124, 65, 67**  
**head starts at 53**
- In that order. If the disk head is initially at cylinder 53:
  - It will first move from 53 to 98, then to 183, 37, 122, 14, 124, 65, and finally to 67.
  - For a total head movement of 640 cylinders.





# First-Come, First-Served (FCFS) (cont.)

- The **wild swing** illustrates the problem with this schedule.
  - The total head movement could be decreased substantially.
  - Performance could be improved.

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53





# Shortest-Seek-Time-First (SSTF)

---


- The SSTF algorithm selects the request with **the least seek time** from the current head position.
- Consider, for example, a disk queue with requests for I/O to blocks on cylinders **98, 183, 37, 122, 14, 124, 65, 67** and head starts at 53
- For our example request queue, **the closest request to the initial head position (53) is at cylinder 65.**

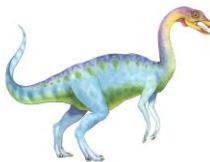
53 → 65 → 67 → 37 → 14 → 98 → 122 → 124



# Shortest-Seek-Time-First (SSTF) (Cont.)

- This scheduling method results in a **total head movement of only 236 cylinders**
- Reduces seek time but may lead to **starvation** for distant requests.






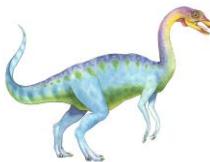
# SCAN Scheduling

- Sometimes called the **elevator algorithm**
- The disk arm **starts at one end** of the disk, and moves toward the other end, servicing requests until it gets to the other end of the disk, where **the head movement is reversed** and servicing continues.
- Total head movement of 208 cylinders.

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53






## SCAN Scheduling

(Cont.)


- Assuming a uniform distribution of requests for cylinders, consider the density of requests when the head reaches one end and reverses direction.
- At this point, relatively few requests are immediately in front of the head, since these cylinders have recently been serviced.
- The heaviest density of requests is at the other end of the disk.
- These requests have also waited the longest, so why not go there first? That is the idea of the next algorithm.



# C-SCAN Scheduling

---

- Circular SCAN (C-SCAN) scheduling is a variant of SCAN designed to provide a more uniform wait time.
- Like SCAN, C-SCAN moves the head from one end of the disk to the other, servicing requests along the way.
- When the head reaches the other end, however, it immediately returns to the beginning of the disk without servicing any requests on the return trip (Total number of cylinders?)




# Assignment.

Consider initial head position 53 and request sequence [98, 183, 37, 122, 14, 124, 65, 67 ] Compute the total distance traveled using C-SCAN Assume Disk size = 0 to 199.

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53






# Assignment. Solution

---

Consider initial head position 53 and request sequence [98, 183, 37, 122, 14, 124, 65, 67 ] Compute the total distance traveled using C-SCAN Assume Disk size = 0 to 199.

**Smart Solution:**

$$= (199-53) + 199 + (37) = 382$$




# LOOK Scheduling

- LOOK and C-LOOK scheduling are versions of SCAN and C-SCAN that look for a request before continuing to move in a given direction.
- The arm goes only as far as the final request in each direction.
- Then, it reverses direction immediately, without going all the way to the end of the disk.

queue = 98, 183, 37, 122, 14, 124, 65, 67

head starts at 53






# NVM Scheduling


---

- NVM Scheduling (Non-Volatile Memory Scheduling) for SSD is fundamentally **different** from traditional disk scheduling due to the lack of mechanical movement.
- Optimizing **parallelism**, **wear leveling**, and **queue management** is critical for modern SSDs and persistent memory.
- **FCFS** is about all that's needed for NVMs, with some variation, like merging adjacent writes.
- NVMs are much **faster** than HDDs at random reads and writes.
- HDDs can sometimes be as fast as NVMs for sequential I/O, especially **sequential writing**.



---

# The End






# Disk Management

---

- **Disk Partitioning** OS allows the disk to be divided into multiple partitions (e.g., C: drive, D: drive).
- **File System Management** OS **formats** disks using file systems like NTFS, FAT32, ext4, etc.
- **Disk Scheduling and Optimization** OS schedules read/write operations efficiently.
- **Disk Space Allocation** OS decides how storage space is allocated to files and directories.
- **Disk Protection and Security** Implements access control and user permissions.
- **Disk Error Handling & Backup** OS detects bad sectors and recovers lost files.



# Logical Formatting (High-Level Formatting)

---

- **Boot Sector:** the first sector of a storage device (hard drive, SSD, USB).  
contains the boot loader, which helps start the operating system. Stores information about the file system type, disk layout, and partition details.
- **File Allocation Tables (FAT):** a data structure that keeps track of where files are stored on the disk. Maps blocks to files.
- **Directory Structures:** organizes files and folders in a hierarchical format.  
Each directory (folder) contains metadata about its files (file name, size, location, (pointers to data blocks), permissions)