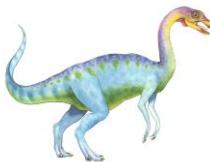


Operating System 2- CS402

Lecture 4 I/O Interface, Kernel I/O Subsystem

2025



Learning Objectives

- Explain **structuring techniques and interfaces** for the operating system that enable I/O devices to be treated in a standard, uniform way.
- Discuss the **main functions** of Kernel I/O Subsystem.
- Explain the **performance aspects** of I/O hardware and software.



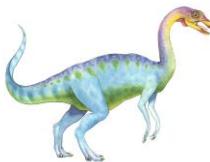
Uniform I/O Device Management

- How an **application can open a file** on a disk without knowing what kind of disk it is?
- How **new disks and other devices can be added** to a computer without disruption of the operating system?




Uniform I/O Device Management (cont.)

- Like other complex software-engineering problems, the approach here involves:
 - **Abstraction**: OS provides a unified interface (e.g., system calls like **open()**, **read()**, **write()**) to interact with devices without knowing hardware details.



A Kernel I/O Structure




How the I/O-related portions of the kernel are structured in software layers.



Characteristics of I/O devices

- Devices vary in many **dimensions**:
- **Data-transfer mode**:
 - A **character-stream** device transfers bytes one by one (keyboard, mouse).
 - A **block** device transfers a block of bytes as a unit (hard drives, SSDs).
- **Access method**:
 - A **sequential** device transfers data in a fixed order defined by the device (magnetic tapes, modem).
 - A **random-access** device user can instruct the device to seek to any of the available data storage locations (hard drives, SSDs).
- **Transfer schedule**:
 - A **synchronous** device performs data transfers with predictable response times (magnetic tapes).
 - An **asynchronous** device exhibits irregular or unpredictable response times (keyboard, mouse).



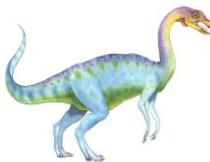
Characteristics of I/O devices

- Devices vary in many dimensions:
 - Sharing:
 - ▶ A **sharable** device can be used concurrently by several processes (hard drives).
 - ▶ A **dedicated** device cannot (keyboard, mouse).
 - Device speed:
 - ▶ Device speeds range from a few bytes per second to a few giga bytes per second.
 - ▶ Slow (keyboard, mouse), Fast (SSD)
 - I/O direction:
 - ▶ Read-write, Read only, or Write only.
 - ▶ Some devices perform both input and output, but others support only one data transfer direction.




Characteristics of I/O devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk



Kernel I/O Subsystem

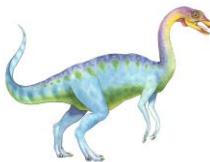


kernel's I/O subsystem provide **services** to efficiently manage
I/O operations:



Kernel I/O Subsystem

- **Scheduling:** Manages the **order and timing** of I/O requests to optimize system performance.
- **Buffering:** Temporarily holds data to accommodate **speed differences** between devices (e.g., keyboard input vs. disk storage).
- **Caching:** Stores frequently **accessed data in fast memory** to speed up retrieval.
- **Spooling:** Queues data for devices that process **one task at a time** (e.g., printer spooling).
- **Device Reservation:** Ensures **exclusive access** to devices when necessary to prevent conflicts.
- **Error Handling:** Detects, reports, and manages I/O errors for system stability.



Buffering in File Reading

- What is Buffering?
 - The process of **temporarily storing data** in memory to reduce direct disk access and improve performance.
 - It allows efficient reading and writing operations by **handling data in chunks instead of individual bytes**.
- Example: Buffered File Reading
 - Reading a large text file efficiently without excessive disk access.
- How It Works:
 1. OS **reads a block of data** (e.g., 4KB or 8KB) from the disk into a buffer in memory.
 2. The program reads data from this buffer instead of making frequent disk requests.
 3. When the **buffer is exhausted**, the OS fetches the next block from the disk.



Caching in the I/O Subsystem

- A cache is a **high-speed memory** that stores frequently accessed data or instructions.
- It helps in **reducing access time** and improving system performance.
- Examples:
 - **Web Caching:** OS and web browsers cache web pages and resources to minimize network traffic and load times.
 - ▶ **Example:** When visiting a website for the second time, cached data speeds up the page load.
 - **Video Caching:** Streaming platforms preload parts of a video to avoid buffering.
 - ▶ **Example:** YouTube loads the next few seconds of a video in advance for smooth playback.




Caching in the I/O Subsystem

- Impact on Performance:
 - Reduces **disk I/O latency** by fetching data from memory instead of slow disks.
 - Improves **system responsiveness** for applications relying on frequent I/O operations.



Introduction to STREAMS

- A **stream** is a full-duplex connection between a device driver and a user-level process.
- Implemented in Unix System V and beyond.
- It consists of :
 - **Stream head**: interfaces with the user process.
 - **Driver end**: controls the device.
 - Zero or more **stream modules** between the stream head and the driver end.





How STREAMS Work

- STREAMS use **message-based communication**.
- Data transferred between queues as **messages**.
- Message Flow in STREAMS: Each module in a STREAM has **two queues**:
 - **Read Queue** → Handles incoming data (device → user).
 - **Write Queue** → Handles outgoing data (user → device).
- Messages are passed through these queues **asynchronously**.



The End

